

Logistic Regression

Outline

- Task Introduction
- Feature extraction
- Code
- Improvement Tips

Task Introduction

- Binary Classification
- Adult Data Set


Attribute, Input



Salary condition, Output

age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hours_per_week	native_country	
39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-manage	Husband	White	Male	0	0	13	United-States	<=50K
38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K
37	Private	284582	Masters	14	Married-civ-spouse	Exec-manage	Wife	White	Female	0	0	40	United-States	<=50K
49	Private	160187	9th	5	Married-spouse	Other-service	Not-in-family	Black	Female	0	0	16	Jamaica	<=50K
52	Self-emp-not-inc	209642	HS-grad	9	Married-civ-spouse	Exec-manage	Husband	White	Male	0	0	45	United-States	>50K
31	Private	45781	Masters	14	Never-married	Prof-specialty	Not-in-family	White	Female	14084	0	50	United-States	>50K
42	Private	159449	Bachelors	13	Married-civ-spouse	Exec-manage	Husband	White	Male	5178	0	40	United-States	>50K
37	Private	280464	Some-college	10	Married-civ-spouse	Exec-manage	Husband	Black	Male	0	0	80	United-States	>50K
30	State-gov	141297	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	Asian-Pac-Islander	Male	0	0	40	India	>50K
23	Private	122272	Bachelors	13	Never-married	Adm-clerical	Own-child	White	Female	0	0	30	United-States	<=50K
32	Private	205019	Assoc-acdm	12	Never-married	Sales	Not-in-family	Black	Male	0	0	50	United-States	<=50K

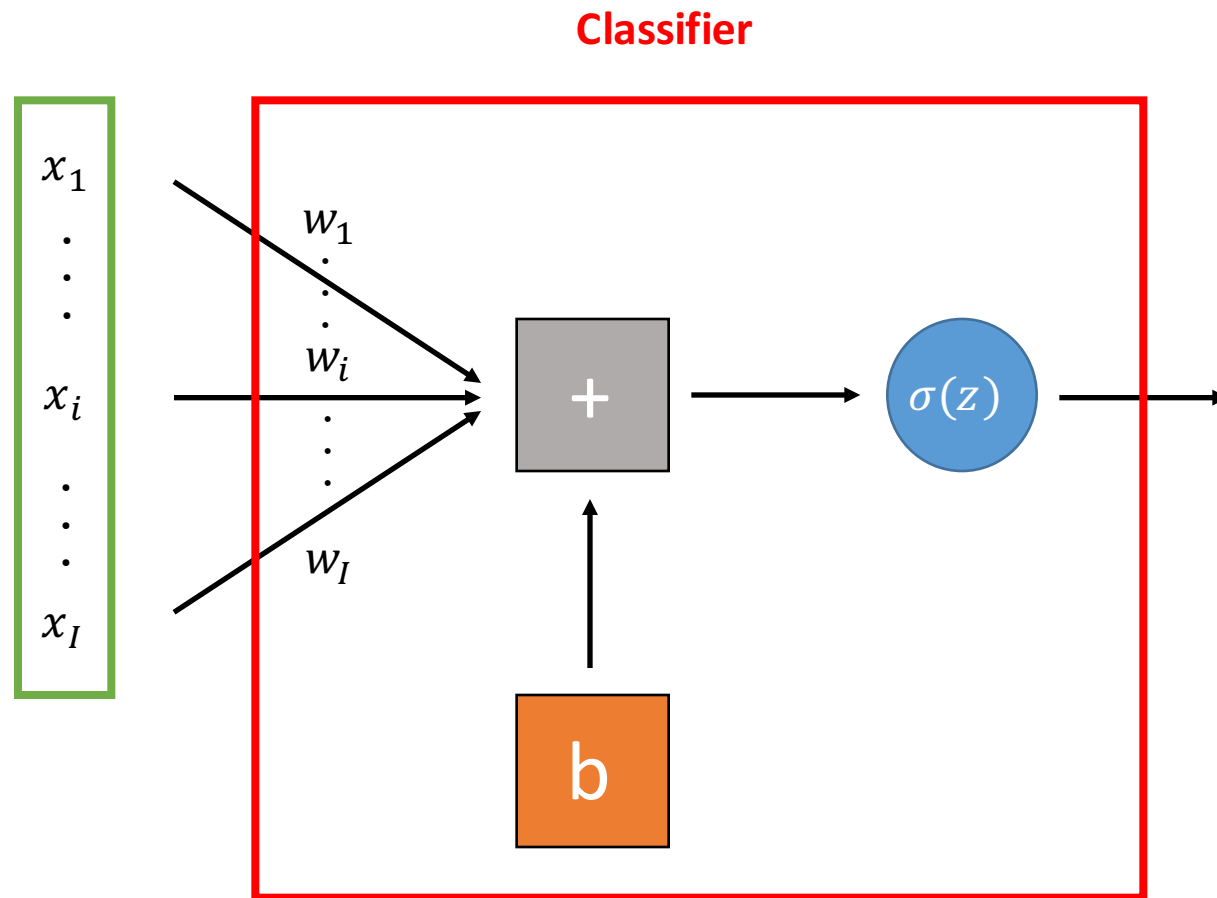
Feature extraction



ex:

1. 15歲, 女, 台灣 : [15, 1, 1, 0]
2. 20歲, 男, 美國 : [20, 0, 0, 1]

[age, female_or_male, is_Taiwan, is_US]



Feature

Feature: One-hot encoding

Label:
0 : <= 50k
1 : > 50K

1	age	fnlwgt	sex	capital_gain	capital_loss	hours_per_week	Federal-gov	Local-gov	Never-worked	Private	Self-
2	39	77516	1	2174	0	40	0	0	0	0	0
3	50	83311	1	0	0	13	0	0	0	0	0
4	38	215646	1	0	0	40	0	0	0	0	0
5	53	234721	1	0	0	40	0	0	0	0	0
6	28	338409	0	0	0	40	0	0	0	0	0
7	37	284582	0	0	0	40	0	0	0	0	0
8	49	160187	0	0	0	16	0	0	0	0	0
9	52	209642	1	0	0	45	0	0	0	0	0
10	31	45781	0	14084	0	50	0	0	0	0	0
11	42	159449	1	5178	0	40	0	0	0	0	0
12	37	280464	1	0	0	80	0	0	0	0	0
13	30	141297	1	0	0	40	0	0	0	0	0
14	23	122272	0	0	0	30	0	0	0	0	0
15	32	205019	1	0	0	50	0	0	0	0	0
16	40	121772	1	0	0	40	0	0	0	0	0
17	34	245487	1	0	0	45	0	0	0	0	0
18	25	176756	1	0	0	35	0	0	0	0	0
19	32	186824	1	0	0	40	0	0	0	0	0
20	38	28887	1	0	0	50	0	0	0	0	0
21	43	292175	0	0	0	45	0	0	0	0	0
22	40	193524	1	0	0	60	0	0	0	0	0
23	54	302146	0	0	0	20	0	0	0	0	0
24	35	76845	1	0	0	40	1	0	0	0	0
25	43	117037	1	0	2042	40	0	0	0	0	0
26	59	109015	0	0	0	40	0	0	0	0	0
27	56	216851	1	0	0	40	0	1	0	0	0
28	19	168294	1	0	0	40	0	0	0	0	0
29	54	180211	1	0	0	60	0	0	0	0	0
30	39	367260	1	0	0	80	0	0	0	0	0
31	49	193366	1	0	0	40	0	0	0	0	0
32	23	190709	1	0	0	52	0	1	0	0	0
33	20	266015	1	0	0	44	0	0	0	0	0
34	45	386940	1	0	1408	40	0	0	0	0	0
N ± master [1:X_train] unix utf-8 no ft 0% 1:1											

1	label
2	0
3	0
4	0
5	0
6	0
7	0
8	0
9	1
10	1
11	1
12	1
13	1
14	0
15	0
16	1
17	0
18	0
19	0
20	0
21	1
22	1
23	0
24	0
25	0
26	0
27	1
28	0
29	1
30	0
31	0
32	0
33	0
34	0
N ± master [1:Y_train]	
[1:Y_train]	

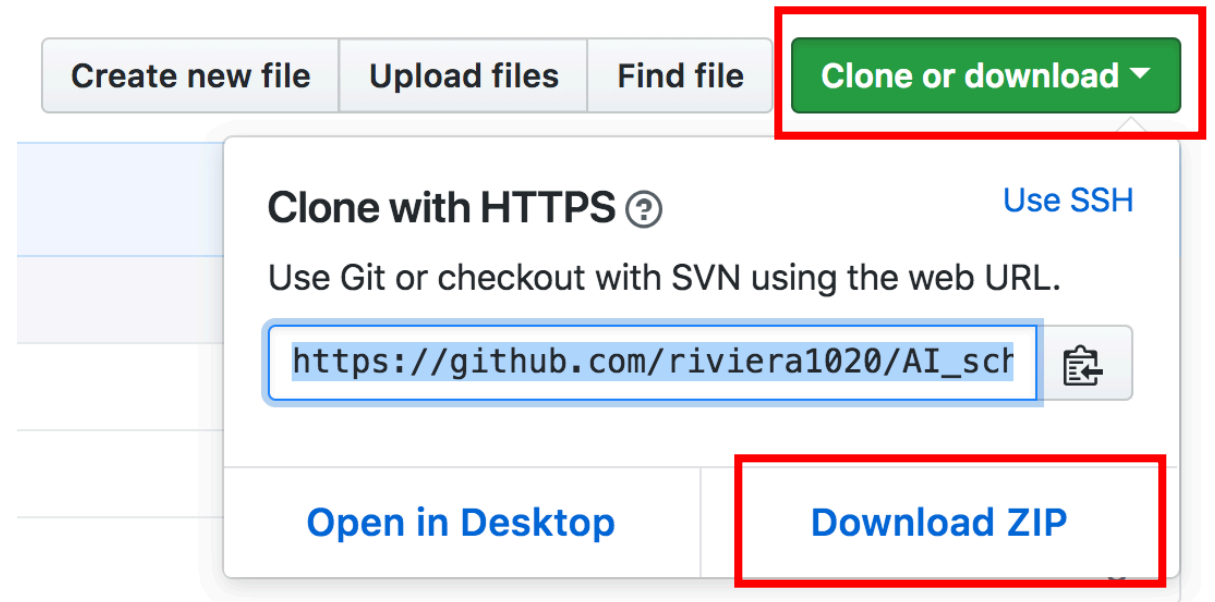
Download

- [Link](#)

- Download zip
- `git clone https://github.com/riviera1020/AI_school_logistic_regression.git`

- File

- `main.py` : 主程式
- `answer.py` : 正確答案
- `feature/` : 幫大家抽好啦
 - `X_train, Y_train` : training data
 - `X_test, ans.csv` : testing data
- `logitsic_params/` : 存model



Code

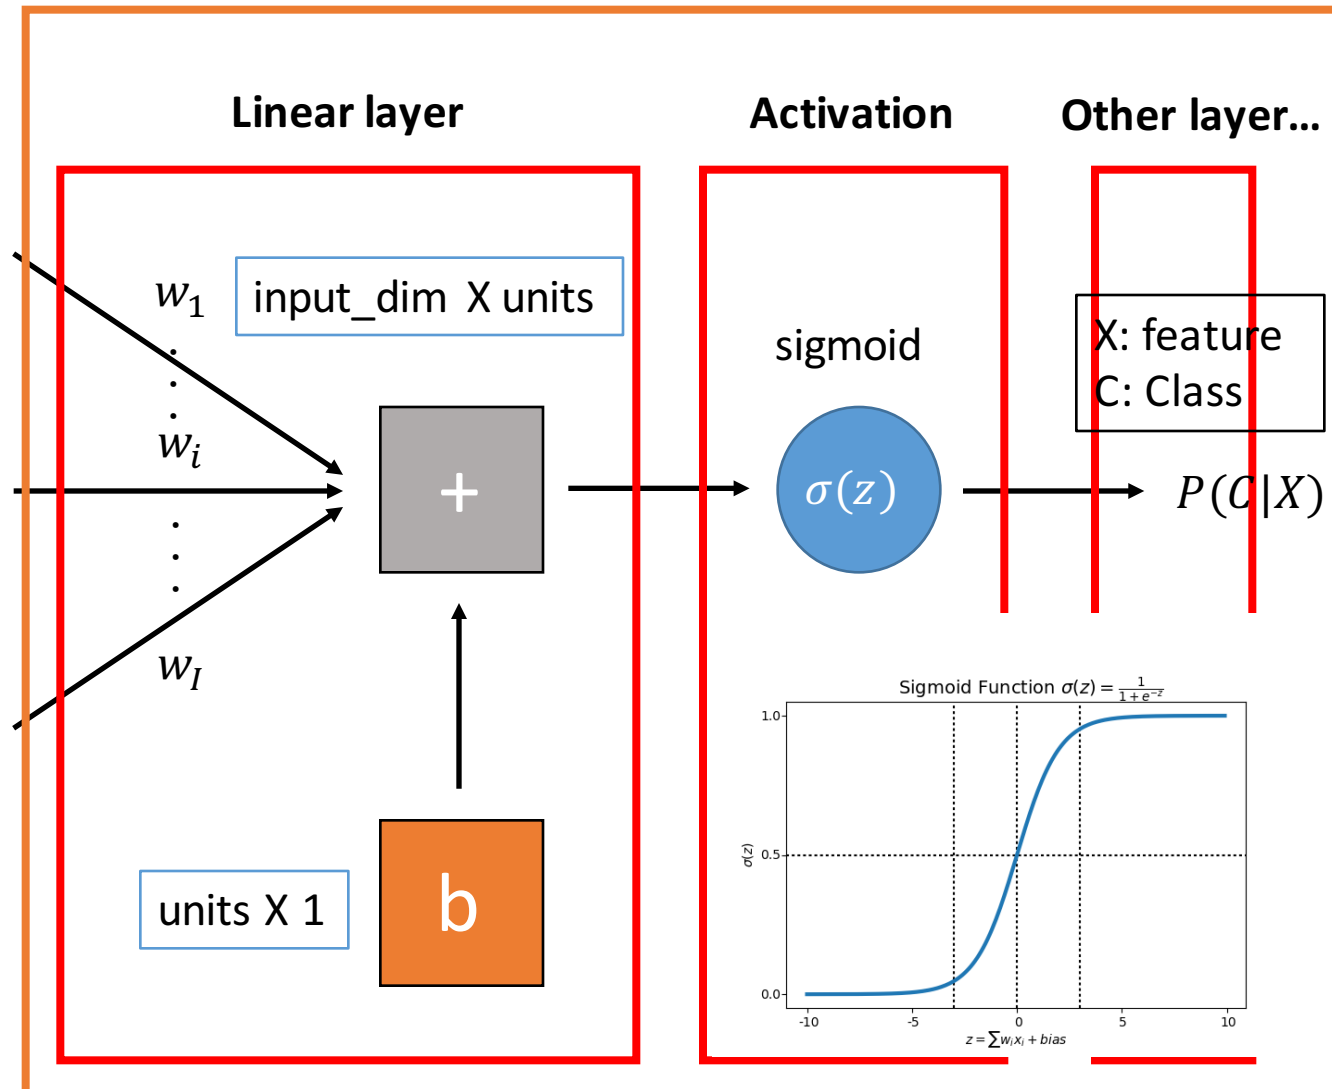
- 用編輯器開main.py
 - Vim, Sublime text, NotePad
- 流程
 - Define model architecture
 - Define model optimizer
 - Compile model
 - Training
 - Inference
- Three step
 - Step 1: Function Set
 - Step 2: Goodness of a Function
 - Step 3: Find the best function

```
def train(X_train, Y_train):  
  
    # TODO, 1  
    # Define model arch  
  
    # TODO, 2  
    # Define optimizer  
  
    # TODO, 3  
    # Compile model  
  
    # TODO, 4  
    # Start training  
  
def infer(X_test, Y_test):  
  
    # TODO, 5  
    # load and inference
```

Define model architecture

Sequential

- 容器：Sequential()
 - `model = Sequential()`
 - `model.add(layer_object)`
- Linear layer
 - `Dense(input_dim, units, use_bias)`
 - `input_dim`: feature的維度
 - `units`: 輸出的維度
 - `use_bias`: True
- Activation
 - `Activation('sigmoid')`



Define model architecture - Code

Index

0

1

X_all.shape : [data_num, feat_dim]

```
# TODO, 1
# Define model arch
model = Sequential()
feat_dims = X_train.shape[1]
model.add(Dense(input_dim = feat_dims, units = 1, use_bias = True))
model.add(Activation('sigmoid'))
```


Step 1: Function Set

Define model optimizer

- Stochastic Gradient Descent

```
# TODO, 2  
# Define optimizer  
sgd = optimizers.SGD(lr=0.001)
```

- Arguments
 - lr : learning rate, 走多大步



Compile model

$$BCE = - \frac{1}{\underbrace{N}_{\text{Batch size}}} \sum_{i=0}^N y_i \cdot \log(\hat{y}_i) + (1 - \underbrace{y_i}_{\text{predict}}) \cdot \log(1 - \underbrace{\hat{y}_i}_{\text{answer}})$$


```
# TODO, 3
# Compile model
model.compile(loss = 'binary_crossentropy',
              optimizer = sgd,
              metrics = ['accuracy'])
```

- Arguments

- loss : cross entropy \longrightarrow Step 2: Goodness of a Function
- optimizer : sgd \longrightarrow Step 3: Find the best function
- metrics : 計算accuracy的方式

Start training

看data幾次



```
# TODO, 4  
# Start training  
model.fit(X_train, Y_train, epochs = 10)  
model.save('./logistic_params/logistic.h5')
```

Training Code

```
def train(X_train, Y_train):  
  
    # TODO, 1  
    # Define model arch  
    model = Sequential()  
    feat_dims = X_train.shape[1]  
    model.add(Dense(input_dim = feat_dims, units = 1, use_bias = True))  
    model.add(Activation('sigmoid'))  
  
    # TODO, 2  
    # Define optimizer  
    sgd = optimizers.SGD(lr=0.001)  
  
    # TODO, 3  
    # Compile model  
    model.compile(loss = 'binary_crossentropy',  
                  optimizer = sgd,  
                  metrics = ['accuracy'])  
  
    # TODO, 4  
    # Start training  
    model.fit(X_train, Y_train, epochs = 10)  
    model.save('./logistic_params/logistic.h5')
```

Inference

```
def infer(X_test, Y_test):  
  
    # TODO, 5  
    # load and inference  
    model = load_model('./logistic_params/logistic.h5')  
    loss, acc = model.evaluate(X_test, Y_test)  
  
    print('Accuracy : ' + str(acc))
```

Run Code

Type commands in command line

Train:

```
python main.py --train
```

Test:

```
python main.py --infer
```

Try to adjust the following parameters:

- epochs
- lr

```
:
name: GeForce GTX 960 major: 5 minor: 2 memoryClockRate(GHz): 1.291
pciBusID: 0000:01:00.0
totalMemory: 3.94GiB freeMemory: 3.25GiB
2018-08-19 17:57:00.091913: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1227] Device peer to peer matrix
2018-08-19 17:57:00.091925: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1233] DMA: 0 1
2018-08-19 17:57:00.091930: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1243] 0: Y N
2018-08-19 17:57:00.091933: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1243] 1: N Y
2018-08-19 17:57:00.091941: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1312] Adding visible gpu devices: 0,
1
2018-08-19 17:57:00.417698: I tensorflow/core/common_runtime/gpu/gpu_device.cc:993] Creating TensorFlow device (/jo
b:localhost/replica:0/task:0/device:GPU:0 with 5641 MB memory) -> physical GPU (device: 0, name: GeForce GTX 1060 6
GB, pci bus id: 0000:05:00.0, compute capability: 6.1)
2018-08-19 17:57:00.437818: I tensorflow/core/common_runtime/gpu/gpu_device.cc:993] Creating TensorFlow device (/jo
b:localhost/replica:0/task:0/device:GPU:1 with 2978 MB memory) -> physical GPU (device: 1, name: GeForce GTX 960, p
ci bus id: 0000:01:00.0, compute capability: 5.2)
32561/32561 [=====] - 2s 62us/step - loss: 0.4722 - acc: 0.7804
Epoch 2/10
32561/32561 [=====] - 1s 41us/step - loss: 0.3488 - acc: 0.8439
Epoch 3/10
32561/32561 [=====] - 1s 40us/step - loss: 0.3345 - acc: 0.8465
Epoch 4/10
32561/32561 [=====] - 1s 41us/step - loss: 0.3291 - acc: 0.8484
Epoch 5/10
32561/32561 [=====] - 1s 40us/step - loss: 0.3261 - acc: 0.8492
Epoch 6/10
32561/32561 [=====] - 1s 40us/step - loss: 0.3241 - acc: 0.8500
Epoch 7/10
32561/32561 [=====] - 1s 40us/step - loss: 0.3228 - acc: 0.8503
Epoch 8/10
32561/32561 [=====] - 1s 41us/step - loss: 0.3217 - acc: 0.8516
Epoch 9/10
32561/32561 [=====] - 1s 40us/step - loss: 0.3211 - acc: 0.8511
Epoch 10/10
32561/32561 [=====] - 1s 41us/step - loss: 0.3204 - acc: 0.8516
riviera1020@531:~/AI_School$ |
```

Improvement tips


- 你其實用過了
 - Normalize
- Overfitting
 - Early stop
- Gradient Descent
 - Momentum
 - Adaptive Learning Rate
- Deep Network
 - Other activation

Normalize

Attribute:
Age, capital_gain \leftrightarrow one-hot


ex:

1. 15歲, 女, 台灣 : [15, 1, 1, 0]
2. 20歲, 男, 美國 : [20, 0, 0, 1]



Source of figure: <http://cs231n.github.io/neural-networks-2/>

Normalize - how



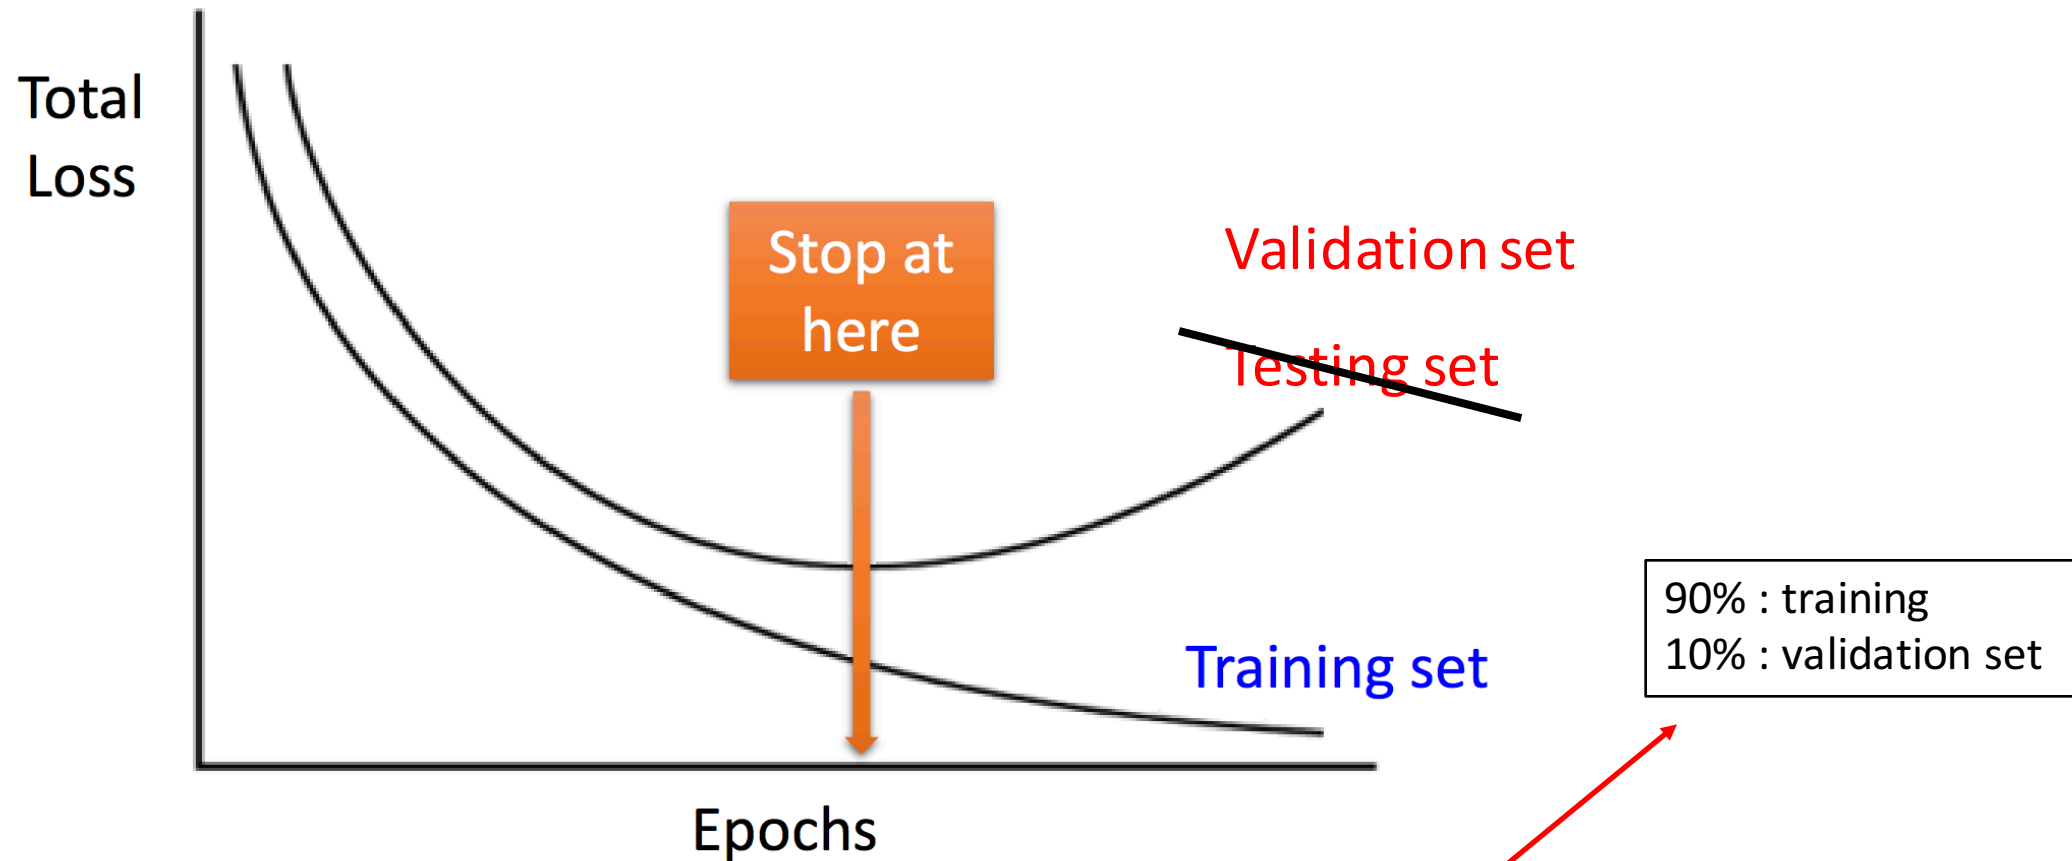
關掉Normalize

```
def main(opts):  
    # Load feature and label  
    X_train, Y_train, X_test, Y_test = load_data(opts.train_data_path,  
                                                  opts.train_label_path,  
                                                  opts.test_data_path,  
                                                  opts.test_label_path)  
  
    # Normalization  
    #X_train, X_test = normalize(X_train, X_test)  
  
    # To train or to infer  
    if opts.train:  
        train(X_train, Y_train)  
    elif opts.infer:  
        infer(X_test, Y_test)  
    else:  
        print("Error: Argument --train or --infer not found")  
    return
```

Comment掉試試看

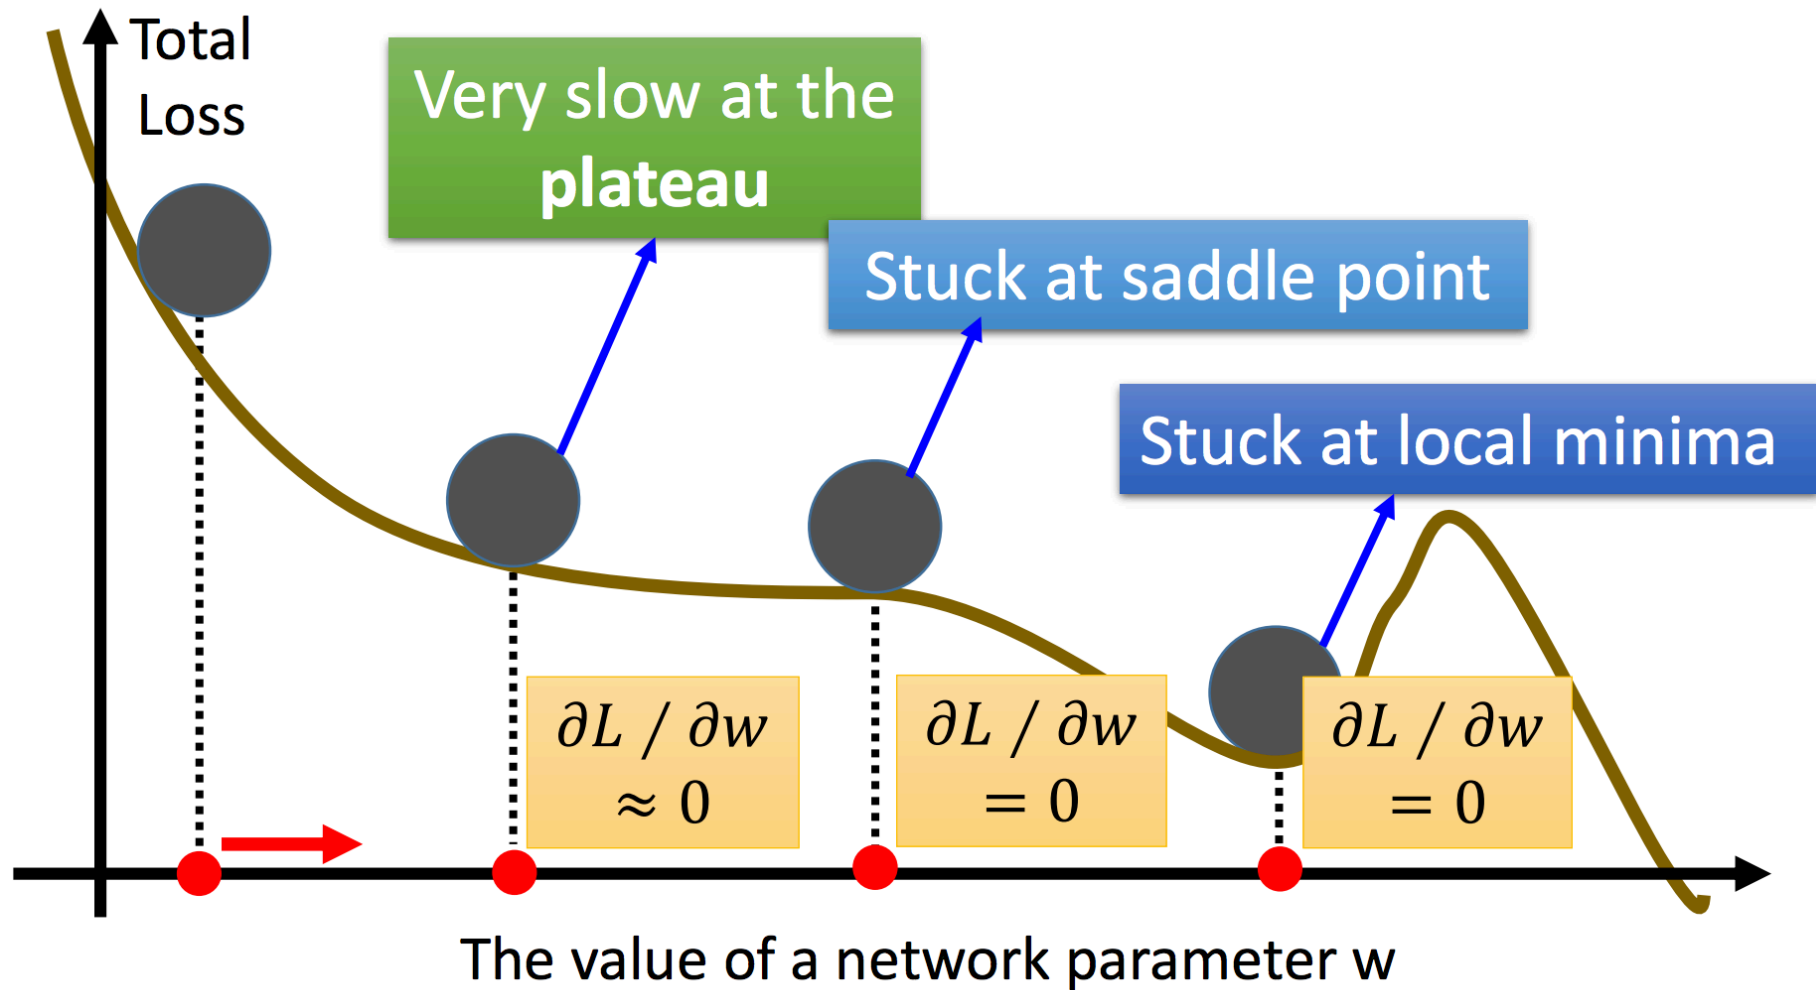
Overfitting – early stop

[Implement Earlystop in keras](#)

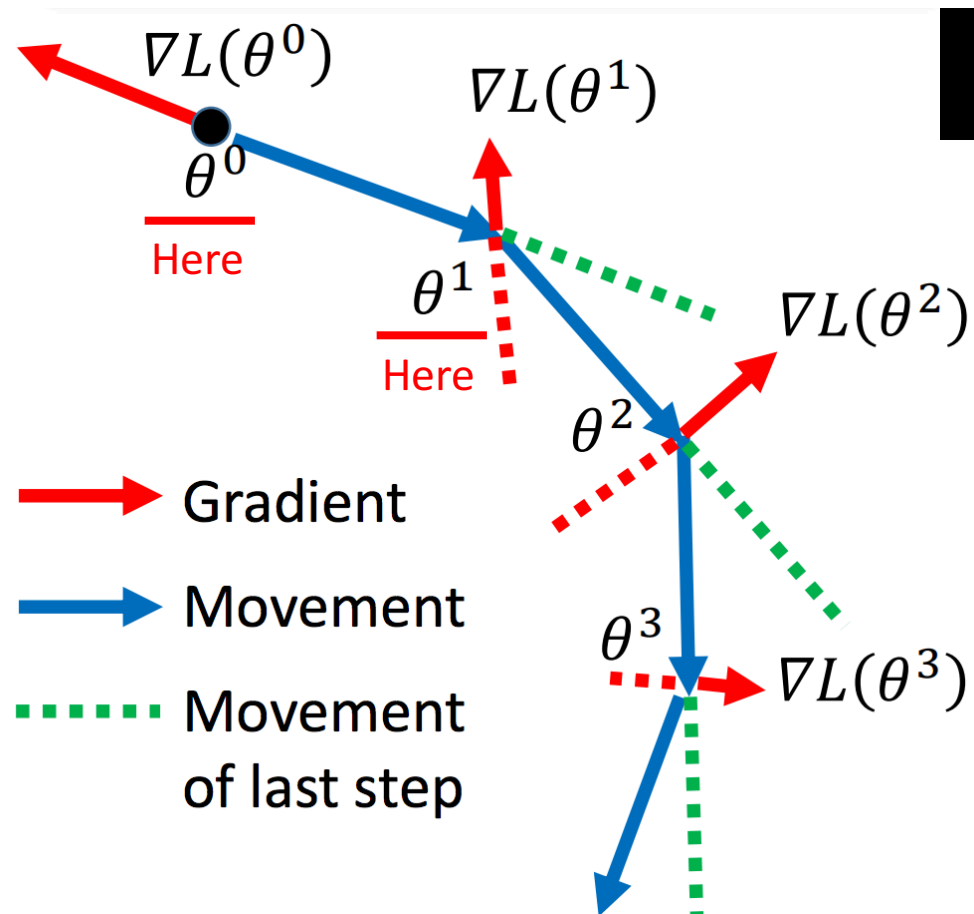


```
model.fit(X_all, Y_all, epochs = 10, validation_split = 0.1)
```

Gradient Descent - momentum



Gradient Descent - momentum



```
sgd = optimizers.SGD(lr=0.001, momentum=0.9)
```

λ
考慮前一次動量的程度

Start at point θ^0

Movement $v^0=0$

Compute gradient at θ^0


Movement $v^1 = \underline{\lambda}v^0 - \eta \nabla L(\theta^0)$

Move to $\theta^1 = \theta^0 + v^1$

Gradient Descent – Adaptive learning rate

- RmsProp

Error Surface can be very complex when training NN.




Gradient Descent – Adaptive learning rate

- Adam: rmsprop + momentum

```
# TODO, 2
# Define optimizer
ada = optimizers.Adam(lr=0.001)

# TODO, 3
# Compile model
model.compile(loss = 'binary_crossentropy',
              optimizer = ada,
              metrics = ['accuracy'])
```



Deep Network



Any dimension you want !

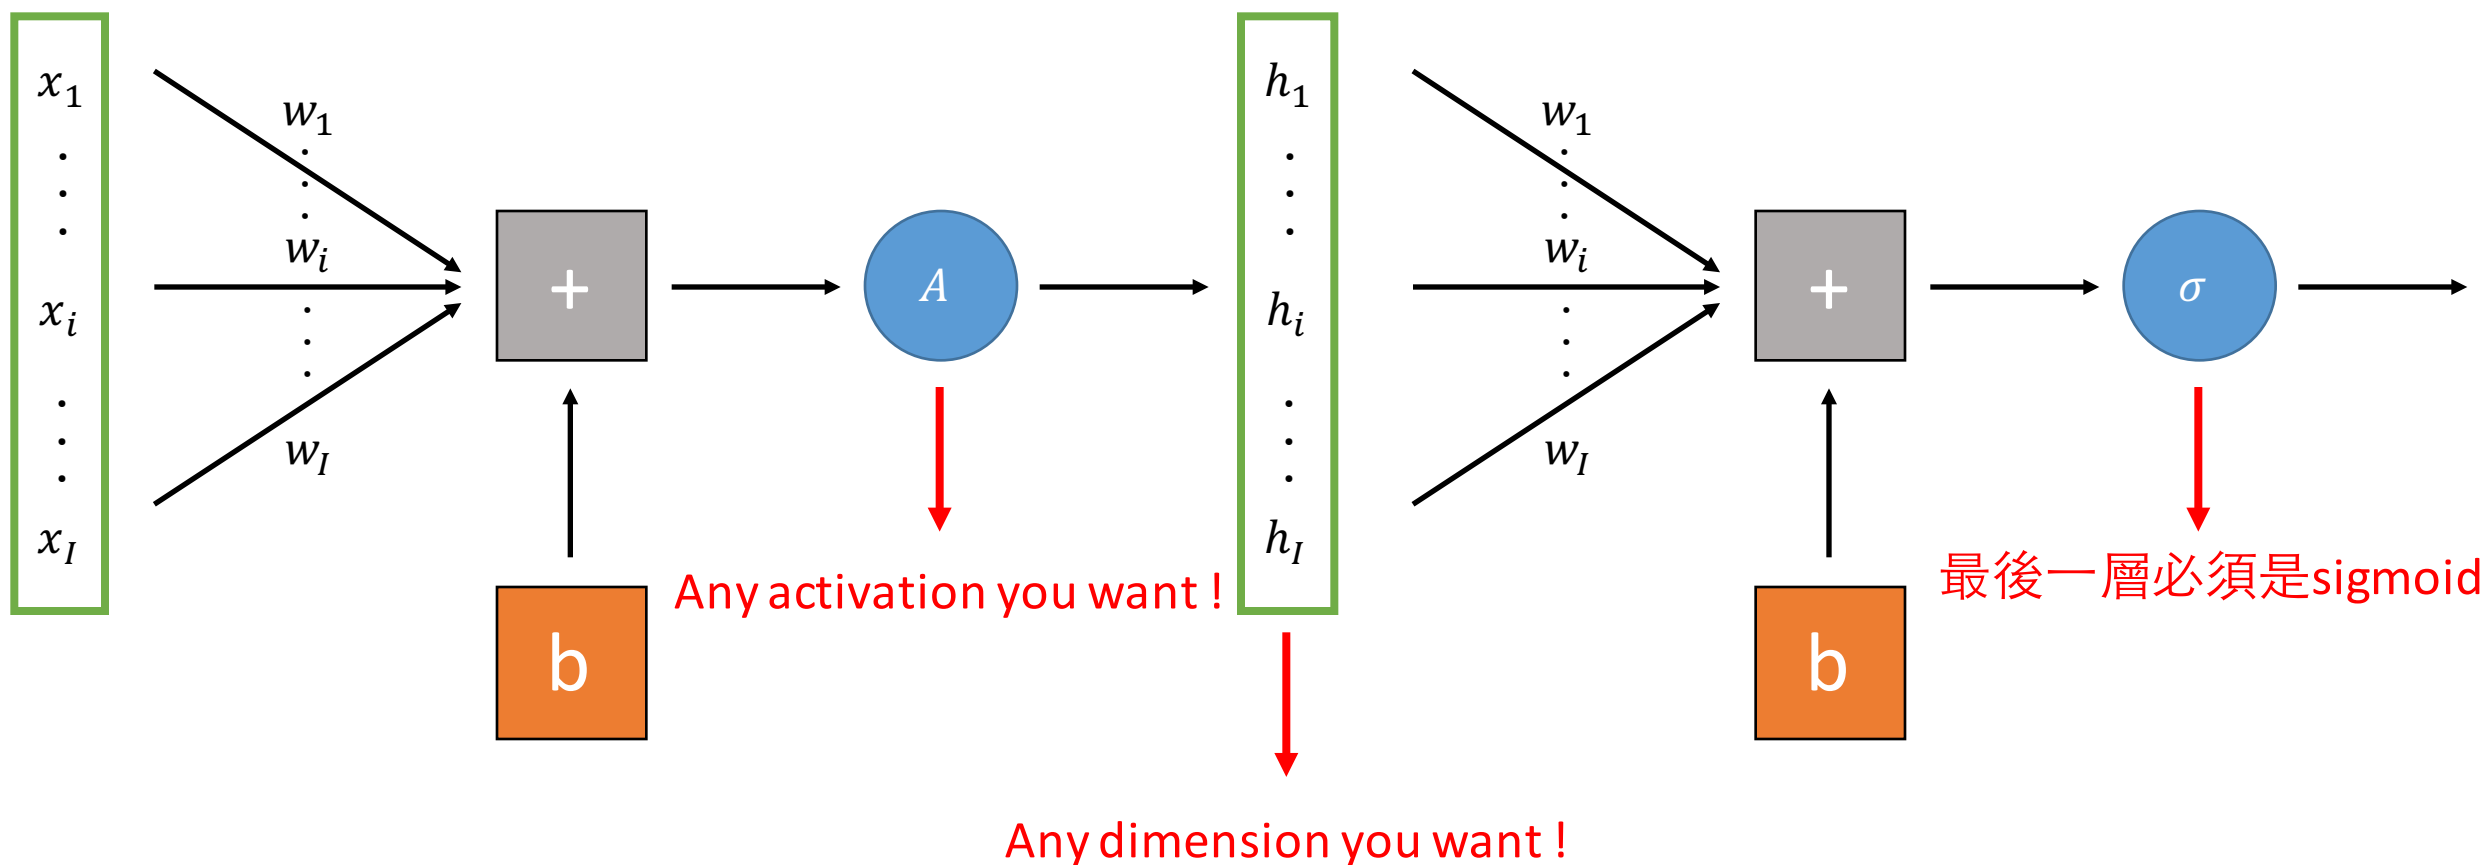
Deep Network

只需要在第一層宣告




```
# TODO, 1
# Define model arch
model = Sequential()
feat_dims = X_train.shape[1]
model.add(Dense(input_dim = feat_dims, units = 1024, use_bias = True))
model.add(Activation('sigmoid'))
model.add(Dense(units = 512, use_bias = True))
model.add(Activation('sigmoid'))
model.add(Dense(units = 1, use_bias = True))
model.add(Activation('sigmoid'))
```

Deep Network



Other activation

- Rectified Linear Unit (ReLU)



```
# TODO, 1
# Define model arch
model = Sequential()
feat_dims = X_train.shape[1]
model.add(Dense(input_dim = feat_dims, units = 1024, use_bias = True))
model.add(Activation('relu'))
model.add(Dense(input_dim = feat_dims, units = 512, use_bias = True))
model.add(Activation('relu'))
model.add(Dense(input_dim = feat_dims, units = 1, use_bias = True))
model.add(Activation('sigmoid'))
```

Implement Time

- Finish original logistic regression
- Adjust epochs, lr
- Try improvement tips

FAQ